

AI BASED PROGRAMMING LAB (CS-634)

MASTER OF TECHNOLOGY

in

COMPUTER SCIENCE AND ENGINEERING

(Artificial Intelligence)

2024-26

Submitted to

Dr. Robin Singh Bhadoria



**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
NATIONAL INSTITUTE OF TECHNOLOGY, HAMIRPUR
NIT, HAMIRPUR
HIMACHAL PRADESH - 177005**

DEC 2024

TABLE OF CONTENTS

PROGRAM NO.	PROGRAM TITLE	PAGE NO.	DATE
1.	Python Program on Data and Variables	3-5	20-Aug-2024
2.	Python Program on Prediction	6-7	27-Aug-2024
3.	Write python programs to implement BFS & DFS Algorithms	8-10	27-Aug-2024
4.	Write a python program to implement Uniform Cost Search	11-12	03-Sept-2024
5.	Write a python program to implement A* Search	13-15	03-Sept-2024
6.	Python Program for Implementing the AO* Algorithm to Solve an And-Or Graph	16-17	10-Sept-2024
7.	Alpha-Beta Pruning for Minimax in Two-Player Game	18	10-Sept-2024
8.	KNN Classification Model for Iris Dataset	19	17-Sept-2024
9.	K-Means Clustering on Customer Data	20-21	17-Sept-2024
10.	Simple Data Operations with NumPy and Pandas	22-23	24-Sept-2024
11.	Text Processing with NLTK: Tokenization, Stopword Removal, and Stemming	24	24-Sept-2024

12.	Implementing a 3-Layer Neural Network for XOR with NumPy	25-27	15-Oct-2024
13.	Deep Learning Model with Binary Classification with Keras for XOR Dataset	28	15-Oct-2024
14.	Setup single node Hadoop cluster and Apply HDFS Commands on Single Node Hadoop Cluster	29-33	22-Oct-2024
15.	Naive Bayes Text Classification with CountVectorizer: Training and Evaluation	34	12-Nov-2024
16.	Python program to calculate Information Gain Calculation for Decision Tree Root Node	35-36	19-Nov-2024

Write a Python program to perform the following operations on a dataset:

1. Load the dataset from a CSV file.
2. Display basic statistics, such as mean, median, standard deviation, and missing values for each column.
3. Handle missing data by either filling it with mean/median values or dropping rows with missing values.
4. Perform data normalization (scaling values between 0 and 1).
5. Group the dataset by a specific column and calculate aggregate statistics (mean, sum, count) for other columns.
6. Visualize the relationships between variables using a pair plot or correlation heatmap.

Use Pandas, NumPy, and Matplotlib/Seaborn libraries to implement the above operations.

Code--

```
import pandas as pd
df = pd.read_csv("/content/drive/MyDrive/annual-enterprise-survey-2023-financial-year-provisional.csv")
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

df.head()

	Year	Industry_aggregation_NZSIOC	Industry_code_NZSIOC	Industry_name_NZSIOC	Units	Variable_code	Variable_name	Variable_c
0	2023	Level 1	99999	All industries	Dollars (millions)	H01	Total income	F perf
1	2023	Level 1	99999	All industries	Dollars (millions)	H04	Sales, government funding, grants and subsidies	F perf
					Dollars		Interest,	F
2	2023	Level 1	99999	All industries	(millions)	H05	dividends and donations	perf
3	2023	Level 1	99999	All industries	Dollars (millions)	H07	Non-operating income	F perf

```
print(df)
```

```

Year Industry_aggregation_NZSIOC Industry_code_NZSIOC \
0      2023                      Level 1          99999
1      2023                      Level 1          99999
2      2023                      Level 1          99999
3      2023                      Level 1          99999
4      2023                      Level 1          99999
...      ...                      ...          ...
50980   2013                      Level 3           ZZ11
50981   2013                      Level 3           ZZ11
50982   2013                      Level 3           ZZ11
50983   2013                      Level 3           ZZ11
50984   2013                      Level 3           ZZ11

Industry_name_NZSIOC          Units Variable_code \
0      All industries Dollars (millions)          H01
1      All industries Dollars (millions)          H04
2      All industries Dollars (millions)          H05
3      All industries Dollars (millions)          H07
4      All industries Dollars (millions)          H08
...      ...                      ...          ...
50980 Food product manufacturing          Percentage          H37
50981 Food product manufacturing          Percentage          H38
50982 Food product manufacturing          Percentage          H39
50983 Food product manufacturing          Percentage          H40
50984 Food product manufacturing          Percentage          H41

Variable_name          Variable_category \
0      Total income          Financial performance
1      Sales, government funding, grants and subsidies          Financial performance
2      Interest, dividends and donations          Financial performance
3      Non-operating income          Financial performance
4      Total expenditure          Financial performance
...      ...                      ...          ...
50980      Quick ratio          Financial ratios
50981      Margin on sales of goods for resale          Financial ratios
50982      Return on equity          Financial ratios
50983      Return on total assets          Financial ratios
50984      Liabilities structure          Financial ratios

Value          Industry_code_ANZSIC06
0      930995 ANZSIC06 divisions A-S (excluding classes K633...
1      821630 ANZSIC06 divisions A-S (excluding classes K633...
2      84354 ANZSIC06 divisions A-S (excluding classes K633...
3      25010 ANZSIC06 divisions A-S (excluding classes K633...
4      832964 ANZSIC06 divisions A-S (excluding classes K633...
...      ...                      ...
50980      52 ANZSIC06 groups C111, C112, C113, C114, C115, ...
50981      40 ANZSIC06 groups C111, C112, C113, C114, C115, ...
50982      12 ANZSIC06 groups C111, C112, C113, C114, C115, ...
50983      5 ANZSIC06 groups C111, C112, C113, C114, C115, ...
50984      46 ANZSIC06 groups C111, C112, C113, C114, C115, ...

[50985 rows x 10 columns]
```

```
df.tail()
```

```

Year Industry_aggregation_NZSIOC Industry_code_NZSIOC Industry_name_NZSIOC Units Variable_code Variable_name Vari
50980 2013                      Level 3          ZZ11      Food product manufacturing Percentage          H37      Quick ratio
50981 2013                      Level 3          ZZ11      Food product manufacturing Percentage          H38      Margin on sales
50982 2013                      Level 3          ZZ11      Food product manufacturing Percentage          H39      of goods for resale
50983 2013                      Level 3          ZZ11      Food product manufacturing Percentage          H39      Return on equity
50984 2013                      Level 3          ZZ11      Food product manufacturing Percentage          H39      equity
```

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50985 entries, 0 to 50984
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Year                                  50985 non-null  int64
1   Industry_aggregation_NZSIOC          50985 non-null  object
2   Industry_code_NZSIOC                 50985 non-null  object
3   Industry_name_NZSIOC                 50985 non-null  object
4   Units                                50985 non-null  object
5   Variable_code                        50985 non-null  object
6   Variable_name                        50985 non-null  object
7   Variable_category                    50985 non-null  object
8   Value                                50985 non-null  object
9   Industry_code_ANZSIC06               50985 non-null  object
dtypes: int64(1),
object(9)memory usage:
3.9+ MB
```

df.describe()

	Year
count	50985.000000
mean	2018.000000
std	3.162309
min	2013.000000
25%	2015.000000
50%	2018.000000
75%	2021.000000
max	2023.000000

df.columns()

```
Index(['Year', 'Industry_aggregation_NZSIOC', 'Industry_code_NZSIOC',
      'Industry_name_NZSIOC', 'Units', 'Variable_code',
      'Variable_name', 'Variable_category', 'Value',
      'Industry_code_ANZSIC06'],
      dtype='object')
```

df.columns.values

```
array(['Year', 'Industry_aggregation_NZSIOC', 'Industry_code_NZSIOC',
      'Industry_name_NZSIOC', 'Units', 'Variable_code',
      'Variable_name', 'Variable_category', 'Value',
      'Industry_code_ANZSIC06'],
      ditypes=object)
```

Write a Python program to predict house prices using a machine learning model.

The program should:

1. Load the housing dataset (CSV file) containing features such as number of rooms, location, area, etc.
2. Preprocess the data by handling missing values, encoding categorical variables, and scaling numeric features.
3. Split the data into training and testing sets.
4. Train a regression model (e.g., Linear Regression, Decision Tree, or SVR) on the training data.
5. Predict the prices of houses from the test set and compare them to the actual prices.
6. Visualize the results by plotting predicted vs. actual prices.

```
from google.colab import drive
drive.mount('/content/drive')
import pandas as pd
df = pd.read_csv("/content/drive/MyDrive/Housing.csv")
```

Mounted at /content/drive

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 545 non-null   int64
1   area                 545 non-null   int64
2   bedrooms             545 non-null   int64
3   bathrooms            545 non-null   int64
4   stories              545 non-null   int64
5   mainroad             545 non-null   object
6   guestroom            545 non-null   object
7   basement             545 non-null   object
8   hotwaterheating      545 non-null   object
9   airconditioning      545 non-null   object
10  parking              545 non-null   int64
11  prefarea             545 non-null   object
12  furnishingstatus     545 non-null   object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
```

df.isnull().sum()

```
0
price      0
area       0
bedrooms   0
bathrooms  0
stories    0
mainroad   0
guestroom  0
basement   0
hotwaterheating  0
airconditioning  0
parking    0
prefarea   0
furnishingstatus  0
```

```
df_area_price=df[['area ','price']]
print(df_area_price)
```

```

↗
   area    price
0   7420  13300000
1   8960  12250000
2   9960  12250000
3   7500  12215000
4   7420  11410000
..    ...    ...
540  3000   1820000
541  2400   1767150
542  3620   1750000
543  2910   1750000
544  3850   1750000

[545 rows x 2 columns]
```

```

from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(df[['area']],df[['price']], test_size =
0.20, random_state = 42)

from sklearn import svm
from sklearn.svm import SVC
model_SVR = svm.SVR()
model_SVR.fit(X_train,Y_train)
Y_pred = model_SVR.predict(X_test) print(X_test,Y_pred)
```

```

↗
   area
316  5900
77   6500
360  4040
90   5000
493  3960
..    ...
15   6000
357  6930
39   6000
54   6000
155  6100
```

```

[109 rows x 1 columns] [4291068.98955744 4291089.5215933 4290965.72765436 4291019.75280652
4290962.19214473 4291093.91286443 4291083.53528354 4291019.14248896
4290941.85393168 4290941.25821675 4291082.48048294 4290961.09167094
4290953.00664623 4290940.51444481 4290962.62186104 4290950.15591473
4290955.63616614 4291073.25297259 4291066.7414097 4291073.25297259
4291020.97246442 4291095.42510977 4290950.72950496 4290955.81963026
4291088.13116148 4291060.37359185 4290940.45163963 4290940.08048196
4291035.35010525 4290940.08048196 4290963.93218271 4290940.65311372
4291073.25297259 4291092.02658269 4291006.02222461 4290995.53258301
4290998.25931194 4290941.17417542 4290946.17130355 4290940.5497816
4291092.24530288 4290950.15591473 4291087.50485529 4290973.42627436
4291094.38265149 4291068.32308962 4291073.25297259 4290997.90219386
4291091.24426482 4290940.08048196 4291097.36479907 4290940.08048196
4291094.88183254 4290989.707461 4290950.45590319 4290940.05254572
4291062.02377763 4290945.74409002 4291088.28758918 4290955.81963026
4290941.17417542 4290951.89952393 4291088.03020756 4291017.31006544
4291041.21305213 4290954.38092232 4291091.7260467 4290940.08048196
4291085.61636812 4291059.55954211 4290940.21945452 4291085.8440915
4291085.61636812 4291054.02930696 4290941.41939442 4290976.50830046
4290950.15591473 4291008.45412197 4291073.25297259 4291037.74064428
4291085.8440915 4290967.57674046 4291049.07050938 4291093.06187824
4290964.91005414 4291071.58560127 4290940.46883351 4290945.49202657
4291097.63268949 4291073.25297259 4290958.45166018 4291091.24426482
4291046.30731748 4290944.64810418 4291082.48048294 4290952.34403498
4291065.72844119 4291064.41864297 4291049.07050938 4291091.7260467
4290941.85393168 4290995.53258301 4290940.03350926 4290942.69336472
4291073.25297259 4291096.57489399 4291073.25297259 4291073.25297259
4291077.19298718]
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1183: DataConversionWarning: A column-vector y was passed when
y = column_or_1d(y, warn=True)
```


Program 3. Write python programs to implement BFS & DFS Algorithms

27-Aug-2024

I. DFS

```
# Maze dimensions and
obstacles maze_size = 6
obstacles = [(0,1),(1,1),(3,2),(3,3),(3,4),(3,5),(0,4),(4,1),(4,2),(4,3)]
start = (0,0)
goal = (0,5)
# checks whether a given position of (x,y) is valid
to move or not def is_valid(x,y):
    return 0 <= x < maze_size and 0 <= y < maze_size and (x,y)
not in obstacles #Dfs function (Depth-first search)
def dfs (current, visited, path):
    x, y = current
    if current == goal:
        path.append(current)
        return True
    visited.add(current)
    moves = [(x-1,y), (x+1, y), (x, y-1), (x, y+1)]
    for move in moves:
        if is_valid(*move) and move not in visited:
            if dfs(move, visited, path):
                path.append(current)
                return True
    return False
#Call DFS function to find the path
visited = set()
path = []
if dfs(start, visited, path):
    path.reverse()
    print("Path found:")
    for position in path:
        print(position)
else:
    print("No path found!")
```

➡ Path found:

```
(0, 0)
(1, 0)
(2, 0)
(3, 0)
(3, 1)
(2, 1)
(2, 2)
(1, 2)
(0, 2)
(0, 3)
(1, 3)
(2, 3)
(2, 4)
(1, 4)
(1, 5)
(0, 5)
```

II. BFS

```
from collections import deque

class GridProblem:
    def __init__(self, initial_state, goal_state, grid):
        # Initializes a grid problem instance with initial and goal states, and the grid layout
        self.initial_state = initial_state
        self.goal_state = goal_state
        self.grid = grid

    def is_goal(self, state):
        # Checks if the given state is the goal state
        return state == self.goal_state

    def is_valid_cell(self, row, col):
        # Checks if the given cell coordinates are within the grid boundaries and not
        blocked
        return 0 <= row < len(self.grid) and 0 <= col < len(self.grid[0]) and
        self.grid[col][row] == 0

    def expand(self, node):
        # Expands the given node by generating child nodes for valid adjacent cells
        row, col = node.state
        children = []
        for dr, dc in [(-1, 0), (1, 0), (0, -1), (0, 1)]:
            new_row, new_col = row + dr, col + dc
            if self.is_valid_cell(new_row, new_col):
                child_state = (new_row, new_col)
                child_node = Node(child_state, parent=node)
                children.append(child_node)
        return children

class Node:
    def __init__(self, state, parent=None, action=None):
        # Initializes a node with a state, parent node (optional), and action (optional)
        self.state = state
        self.parent = parent
        self.action = action

def breadth_first_search(problem):
    # Performs breadth-first search algorithm to find a solution for the given problem
    node = Node(problem.initial_state)
    if problem.is_goal(node.state):
        return node

    frontier = deque([node])
    reached = {problem.initial_state}

    while frontier:
        node = frontier.popleft()

        for child in problem.expand(node):
            state = child.state

            if problem.is_goal(state):
                return child
            if state not in reached:
                reached.add(state)
                frontier.append(child)
    return None

def reconstruct_path(node):
    # Reconstructs the path from the goal node back to the initial node
    path = []
    while node:
        path.append(node.state)
```

```

        node = node.parent
    return list(reversed(path))

def print_complete_path(path):
    # Prints the complete path from start to goal
    if path:
        for step, point in enumerate(path):
            print("Step {}: {}".format(step, point))
    else:
        print("No solution found")

# Example usage and grid definition
"""
    1 : Denotes the obstacles
    0 : Empty space or a non-obstacle cell in the grid
"""
grid = [
    [0, 1, 0, 0, 1, 0, 0],
    [0, 1, 0, 0, 1, 0, 0],
    [0, 0, 0, 0, 1, 0, 0],
    [0, 0, 1, 0, 1, 0, 0],
    [0, 0, 1, 0, 0, 0, 0],
    [0, 0, 1, 0, 0, 0, 0],
    [0, 0, 1, 0, 0, 0, 0]
]

# Define initial and goal states
initial_state = (0, 0)
goal_state = (6, 0)

# Define the problem instance
problem = GridProblem(initial_state, goal_state, grid)

# Perform breadth-first search to find a solution
solution_node = breadth_first_search(problem)

# Print solution if found
print('!! Reached the Goal!!' if solution_node else None)
if solution_node:
    print("Solution found!")
    solution_path = reconstruct_path(solution_node)
    print("Complete Path:")
    print_complete_path(solution_path)
else:
    print("No solution found")

```

```

→ !! Reached the Goal!!
Solution found!
Complete Path:
Step 0: (0, 0)
Step 1: (0, 1)
Step 2: (0, 2)
Step 3: (1, 2)
Step 4: (2, 2)
Step 5: (3, 2)
Step 6: (3, 3)
Step 7: (3, 4)
Step 8: (4, 4)
Step 9: (5, 4)
Step 10: (6, 4)
Step 11: (6, 3)
Step 12: (6, 2)
Step 13: (6, 1)
Step 14: (6, 0)

```

Program 4. Write a python program to implement Uniform Cost Search.

03-Sept-2024

UNIFORM COST SEARCH

```
# Python3 implementation of above approach
def uniform_cost_search(goal, start):
    # minimum cost upto
    # goal state from starting
    global graph, cost
    answer = []
    # create a priority queue
    queue = []
    # set the answer vector to max value
    for i in range(len(goal)):
        answer.append(10**8)
    # insert the starting index
    queue.append([0, start])
    # map to store visited node
    visited = { }
    # count
    count = 0
    # while the queue is not empty
    while (len(queue) > 0):
        # get the top element of the
        queue = sorted(queue)
        p = queue[-1]
        # pop the element
        del queue[-1]
        # get the original value
        p[0] *= -1
        # check if the element is part of
        # the goal list
        if (p[1] in goal):
            index = goal.index(p[1])
            # if a new goal is reached
            if (answer[index] == 10**8):
                count += 1
            # if the cost is less
            if (answer[index] > p[0]):
                answer[index] = p[0]
            # pop the element
            del queue[-1]
            queue = sorted(queue)
            if (count == len(goal)):
                return answer
        # check for the non visited nodes
        # which are adjacent to present node
        if (p[1] not in visited):
            for i in range(len(graph[p[1]])):
                # value is multiplied by -1 so that
                # least priority is at the top
                queue.append([(p[0] + cost[(p[1],
graph[p[1]][i]))]* -1, graph[p[1]][i])]
            # mark as visited
            visited[p[1]] = 1
    return answer
# main function
if __name__ == '__main__':
    # create the graph
```

```

graph,cost = [[] for i in range(8)],{}
# add edge
graph[0].append(1)
graph[0].append(3)
graph[3].append(1)
graph[3].append(6)
graph[3].append(4)
graph[1].append(6)
graph[4].append(2)
graph[4].append(5)
graph[2].append(1)
graph[5].append(2)
graph[5].append(6)
graph[6].append(4)
# add the cost
cost[(0, 1)] = 2
cost[(0, 3)] = 5
cost[(1, 6)] = 1
cost[(3, 1)] = 5
cost[(3, 6)] = 6
cost[(3, 4)] = 2
cost[(2, 1)] = 4
cost[(4, 2)] = 4
cost[(4, 5)] = 3
cost[(5, 2)] = 6
cost[(5, 6)] = 3
cost[(6, 4)] = 7
# goal state
goal = []
# set the goal
# there can be multiple goal states
goal.append(6)
# get the answer
answer = uniform_cost_search(goal, 0)
# print the answer
print("Minimum cost from 0 to 6 is = ",answer[0])

```

➦ Minimum cost from 0 to 6 is = 3

A* SEARCH

```

# Python program for A* Search Algorithm
import math
import heapq
# Define the Cell class
class Cell:
    def __init__(self):

        # Parent cell's row index
        self.parent_i = 0
        # Parent cell's column index
        self.parent_j = 0
    # Total cost of the cell (g + h)
    self.f = float('inf')
    # Cost from start to this cell
    self.g = float('inf')
    # Heuristic cost from this cell to destination
    self.h = 0
# Define the size of the grid
ROW = 9
COL = 10
# Check if a cell is valid (within the grid)
def is_valid(row, col):
    return (row >= 0) and (row < ROW) and (col >= 0) and (col < COL)
# Check if a cell is unblocked
def is_unblocked(grid, row, col):
    return grid[row][col] == 1
# Check if a cell is the destination
def is_destination(row, col, dest):
    return row == dest[0] and col == dest[1]
# Calculate the heuristic value of a cell (Euclidean distance to destination)
def calculate_h_value(row, col, dest):
    return ((row - dest[0]) ** 2 + (col - dest[1]) ** 2) ** 0.5
# Trace the path from source to destination
def trace_path(cell_details, dest):
    print("The Path is ")
    path = []
    row = dest[0]
    col = dest[1]
    # Trace the path from destination to source using parent cells
    while not (cell_details[row][col].parent_i == row and cell_details[row][col].parent_j == col):
        path.append((row, col))
        temp_row = cell_details[row][col].parent_i
        temp_col = cell_details[row][col].parent_j
        row = temp_row
        col = temp_col
    # Add the source cell to the path
    path.append((row, col))
    # Reverse the path to get the path from source to destination
    path.reverse()
    # Print the path
    for i in path:
        print("-", i, end=" ")
    print()
    # Implement the A* search algorithm
def a_star_search(grid, src, dest):
    # Check if the source and destination are valid
    if not is_valid(src[0], src[1]) or not is_valid(dest[0], dest[1]):
        print("Source or destination is invalid")
        return
    # Check if the source and destination are unblocked
    if not is_unblocked(grid, src[0], src[1]) or not is_unblocked(grid, dest[0], dest[1]):

```

```

print("Source or the destination is blocked")
return
# Check if we are already at the destination
if is_destination(src[0], src[1], dest):
    print("We are already at the destination")
    return
# Initialize the closed list (visited cells)
closed_list = [[False for _ in range(COL)] for _ in range(ROW)]
# Initialize the details of each cell
cell_details = [[Cell() for _ in range(COL)] for _ in range(ROW)]
# Initialize the start cell details
i = src[0]
j = src[1]
cell_details[i][j].f = 0
cell_details[i][j].g = 0
cell_details[i][j].h = 0
cell_details[i][j].parent_i = i
cell_details[i][j].parent_j = j
# Initialize the open list (cells to be visited) with the start cell
open_list = []
heapq.heappush(open_list, (0.0, i, j))
# Initialize the flag for whether destination is found
found_dest = False
# Main loop of A* search algorithm
while len(open_list) > 0:
    # Pop the cell with the smallest f value from the open list
    p = heapq.heappop(open_list)
    # Mark the cell as visited
    i = p[1]
    j = p[2]
    closed_list[i][j] = True
    # For each direction, check the successors
    directions = [(0, 1), (0, -1), (1, 0), (-1, 0), (1, 1), (1, -1), (-1, 1), (-1, -1)]
    for dir in directions:
        new_i = i + dir[0]
        new_j = j + dir[1]
        # If the successor is valid, unblocked, and not visited
        if is_valid(new_i, new_j) and is_unblocked(grid, new_i, new_j) and not closed_list[new_i][new_j]:
            # If the successor is the destination
            if is_destination(new_i, new_j, dest):
                # Set the parent of the destination cell
                cell_details[new_i][new_j].parent_i = i
                cell_details[new_i][new_j].parent_j = j
                print("The destination cell is found")
                # Trace and print the path from source to destination
                trace_path(cell_details, dest)
                found_dest = True
                return
            else:
                # Calculate the new f, g, and h values
                g_new = cell_details[i][j].g + 1.0
                h_new = calculate_h_value(new_i, new_j, dest)
                f_new = g_new + h_new
                # If the cell is not in the open list or the new f value is smaller
                if cell_details[new_i][new_j].f == float('inf') or cell_details[new_i][new_j].f > f_new:
                    # Add the cell to the open list
                    heapq.heappush(open_list, (f_new, new_i, new_j))
                    # Update the cell details
                    cell_details[new_i][new_j].f = f_new
                    cell_details[new_i][new_j].g = g_new
                    cell_details[new_i][new_j].h = h_new
                    cell_details[new_i][new_j].parent_i = i
                    cell_details[new_i][new_j].parent_j = j
    # If the destination is not found after visiting all cells
    if not found_dest:
        print("Failed to find the destination cell")
# Driver Code
def main():
    # Define the grid (1 for unblocked, 0 for blocked)
    grid = [
        [1, 0, 1, 1, 1, 1, 0, 1, 1, 1],
        [1, 1, 1, 0, 1, 1, 1, 0, 1, 1],
        [1, 1, 1, 0, 1, 1, 0, 1, 0, 1],
        [0, 0, 1, 0, 1, 0, 0, 0, 0, 1],
        [1, 1, 1, 0, 1, 1, 1, 0, 1, 0],
        [1, 0, 1, 1, 1, 1, 0, 1, 0, 0],
        [1, 0, 0, 0, 0, 1, 0, 0, 0, 1],
        [1, 0, 1, 1, 1, 1, 0, 1, 1, 1],
        [1, 1, 1, 0, 0, 0, 1, 0, 0, 1]
    ]

```

```
# Define the source and destination
src = [8, 0]
dest = [0, 0]
# Run the A* search algorithm a_star_search(grid, src, dest)
if __name__ == "__main__":
    main()
```

The destination cell is found The Path is -> (8, 0) -> (7, 0) -> (6, 0) -> (5, 0) -
> (4, 1) -> (3, 2) -> (2, 1) -> (1, 0) -> (0, 0)

Program 6. Write a Python program to implement the A0* algorithm for solving an And-Or graph. The program should take an And-Or graph as input, search for the optimal solution

10-Sept-2024

```
# Cost to find the AND and
OR path def Cost(H,
condition, weight = 1):
    cost = {}
    if 'AND' in condition:
        AND_nodes = condition['AND']
        Path_A = ' AND '.join(AND_nodes)
        PathA = sum(H[node]+weight for node in
        AND_nodes)cost[Path_A] = PathA
    if 'OR' in condition:
        OR_nodes = condition['OR']
        Path_B = ' OR '.join(OR_nodes)
        PathB = min(H[node]+weight for node in
        OR_nodes)cost[Path_B] = PathB
    return cost
# Update the cost
def update_cost(H, Conditions, weight=1):
    Main_nodes =
    list(Conditions.keys())
    Main_nodes.reverse()
    least_cost= {}
    for key in Main_nodes:
        condition = Conditions[key]
        print(key,':', Conditions[key], '>>>', Cost(H, condition,
        weight))c = Cost(H, condition, weight)
        H[key] = min(c.values())
        least_cost[key] = Cost(H, condition,
        weight)return least_cost

# Print the shortest path
def shortest_path(Start,Updated_cost, H):
    Path = Start
    if Start in Updated_cost.keys():
        Min_cost =
        min(Updated_cost[Start].values())key
        = list(Updated_cost[Start].keys())
        values =
        list(Updated_cost[Start].values())
        Index = values.index(Min_cost)
        # FIND MINIMUM PATH KEY
        Next =
        key[Index].split(
        )# ADD TO PATH
        FOR OR PATH
        if len(Next) == 1:
            Start =Next[0]
            Path += '<--' +shortest_path(Start,
            Updated_cost, H)# ADD TO PATH FOR AND PATH
        else:
            Path += '<--
            ('+key[Index]+' )
            Start = Next[0]
            Path += '[' +shortest_path(Start, Updated_cost, H)
            + ' + 'Start = Next[-1]
            Path += shortest_path(Start, Updated_cost,
            H) + ']'return Path

H = {'A': -1, 'B': 5, 'C': 2, 'D': 4, 'E': 7, 'F': 9, 'G': 3, 'H': 0, 'I':0, 'J':0}
Conditions = {
'A': {'OR': ['B'], 'AND':
['C', 'D']}, 'B': {'OR': ['E',
'F']},
'C': {'OR': ['G'], 'AND':
['H', 'I']}, 'D': {'OR': ['J']}
}
# weight
weight = 1
# Updated cost
print('Updated Cost :')
Updated_cost = update_cost(H, Conditions,
```

```

weight=1)print('*'*75)
print('Shortest Path :\n',shortest_path('A', Updated_cost,H))

⇒ Updated Cost :
D : {'OR': ['J']} >>> {'J': 1}
C : {'OR': ['G'], 'AND': ['H', 'I']} >>> {'H AND I': 2,
'G': 4}B : {'OR': ['E', 'F']} >>> {'E OR F': 8}
A : {'OR': ['B'], 'AND': ['C', 'D']} >>> {'C AND D': 5, 'B': 9}
*****
*****Shortest Path :
A<--(C AND D) [C<--(H AND I) [H + I] + D<--J]

```

Program 7. Write a Python program to implement the Alpha-Beta Pruning algorithm for optimizing the Minimax decision-making process in a two-player game. 10-Sept-2024

```
# Python program to demonstrate working of Alpha-Beta Pruning # Initial values of Alpha and Beta
MAX, MIN = 1000, -1000
# Returns optimal value for current player #(Initially called for root and maximizer)
def minimax(depth, nodeIndex, maximizingPlayer, values, alpha, beta):
    # Terminating condition. i.e # leaf node is reached
    if depth == 3:
        return values[nodeIndex]
    if maximizingPlayer:
        best = MIN
        # Recur for left and right children
        for i in range(0, 2):
            val = minimax(depth + 1, nodeIndex * 2 + i, False, values, alpha, beta)
            best = max(best, val)
            alpha = max(alpha, best) # Alpha Beta Pruning
            if beta <= alpha:
                break
        return best
    else:
        best = MAX
        # Recur for left and # right children
        for i in range(0, 2):
            val = minimax(depth + 1, nodeIndex * 2 + i, True, values, alpha, beta)
            best = min(best, val)
            beta = min(beta, best)
            # Alpha Beta Pruning
            if beta <= alpha:
                break
        return best


# Driver Code

if __name__ == "__main__":
    values = [3, 5, 6, 9, 1, 2, 0, -1]
    print("The optimal value is :", minimax(0, 0, True, values, MIN, MAX))

The optimal value is : 5
```

Program 8. Implement a K-Nearest Neighbours (KNN) classification model using the Iris dataset to predict the species of iris flowers. 17-Sept-2024

```
# Import necessary modules
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection
import train_test_split
from sklearn.datasets import load_iris
from sklearn.metrics import f1_score,
recall_score, accuracy_score# Loading data
irisData = load_iris()
# Create feature and
target arraysX =
irisData.data
y = irisData.target
# Split into training and test set
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)# Initialize KNN classifier
knn =
KNeighborsClassifier(n_ne
ighbors=7)# Fit the model
knn.fit(X_train, y_train)
# Predict on dataset which model
has not seen beforey_pred =
knn.predict(X_test)
# Print predictions
print("Predictions:", y_pred)
# Calculate and print F1 score,
recall, and accuracyf1 =
f1_score(y_test, y_pred,
average='weighted')
recall = recall_score(y_test, y_pred,
average='weighted')accuracy =
accuracy_score(y_test, y_pred)
print(f"F1
Score: {f1}")
print(f"Recall
: {recall}")
print(f"Accuracy: {accuracy}")
```

 Predictions: [1 0 2 1 1 0 1 2 2 1 2 0 0 0 1 2 1 1 2 0 2 0 2 2 2 2 0 0]
F1 Score: 0.9664109121909632
Recall: 0.9666666666666667
Accuracy: 0.9666666666666667

Program 9. Apply the K-Means clustering algorithm on a customer dataset to group customers according to their annual income and spending score.

17-Sept-2024

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

# Load the customer dataset
data = pd.read_csv('Mall_Customers.csv')
data

data.columns.tolist()

X = data[['Annual Income (k$)', 'Spending Score (1-100)']].values


wcss = [] # Within-cluster sum of squares_24_mcs_105
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300, n_init=10, random_state=42)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)

# Plot the Elbow method graph
plt.figure(figsize=(8,6))
plt.plot(range(1, 11), wcss, marker='o', linestyle='--')
plt.title('The Elbow Method')
plt.xlabel('Number of Clusters')
plt.ylabel('WCSS')
plt.grid(True)
plt.show()

optimal_clusters = 5
kmeans = KMeans(n_clusters=optimal_clusters, init='k-means++', max_iter=300, n_init=10, random_state=42)
y_kmeans = kmeans.fit_predict(X)

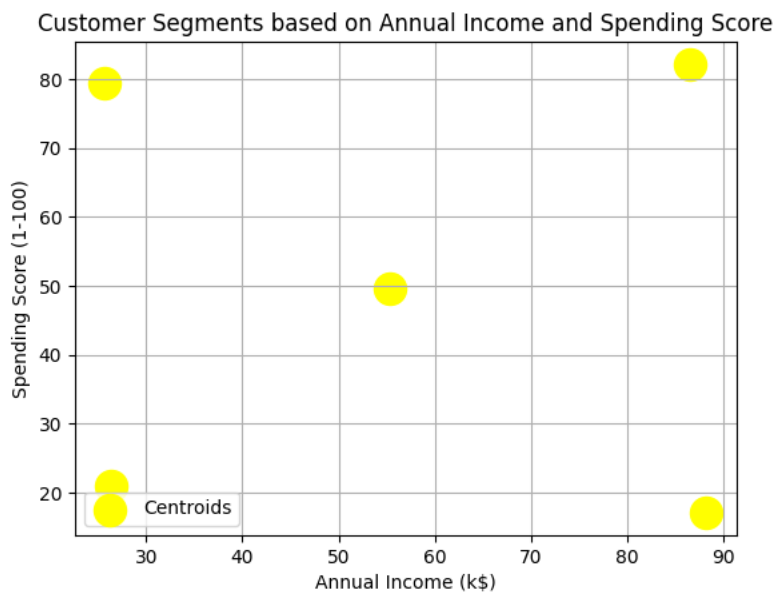


plt.figure(figsize=(8,6))
plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s=100, c='red', label='Cluster 1')
plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s=100, c='blue', label='Cluster 2')
plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s=100, c='green', label='Cluster 3')
plt.scatter(X[y_kmeans == 3, 0], X[y_kmeans == 3, 1], s=100, c='cyan', label='Cluster 4')
plt.scatter(X[y_kmeans == 4, 0], X[y_kmeans == 4, 1], s=100, c='magenta', label='Cluster 5')

plt.scatter(kmeans.cluster_centers_[0, 0], kmeans.cluster_centers_[0, 1], s=300, c='yellow', label='Centroids')
plt.title('Customer Segments based on Annual Income and Spending Score')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()
plt.grid(True)
plt.show()
```

 Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).


	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score
(1-100)0	1	Male	19	15	39
	1	2	Male	21	81
	2	3	Female	20	6
	3	4	Female	23	77
	4	5	Female	31	40

```
['CustomerID', 'Gender', 'Age', 'Annual Income (k$)', 'Spending Score (1-100)']
```



Program 10. an example of a simple Python program that uses both NumPy and Pandas to generate a dataset, create a DataFrame, perform basic operations like adding a new column, sorting, filtering, and calculating statistics?
24-Sept-2024

```
import numpy as np
import pandas as pd
# Step 1: Create random data using NumPy
np.random.seed(42) # For reproducibility
data = np.random.randint(10, 100, size=(10, 3)) # 10 rows and 3 columns
# Step 2: Create a DataFrame using Pandas
df = pd.DataFrame(data, columns=['Math', 'Science', 'English'])
# Display the initial DataFrame
print("Initial DataFrame:")
print(df)
# Step 3: Add a new column for the average score of each student
df['Average'] = df.mean(axis=1)
# Display the DataFrame with the new column
print("\nDataFrame with Average scores:")
print(df)
# Step 4: Sort the DataFrame based on the Average column
df_sorted = df.sort_values(by='Average', ascending=False)
print("\nDataFrame sorted by Average scores (descending order):")
print(df_sorted)
# Step 5: Filter rows where Average score is greater than 50
df_filtered = df[df['Average'] > 50]
print("\nFiltered DataFrame with Average score > 50:")
print(df_filtered)
# Step 6: Calculate basic statistics for the 'Math' column
math_mean = df['Math'].mean()
math_std = df['Math'].std()
print(f"\nStatistics for Math scores:\nMean: {math_mean:.2f}, Standard Deviation: {math_std:.2f}")
```

 Initial DataFrame:

	Math	Science	English
0	61	24	81
1	70	30	92
2	96	84	84
3	97	33	12
4	31	62	11
5	97	39	47
6	11	73	69
7	30	42	85
8	67	31	98
9	58	68	51

DataFrame with Average scores:

	Math	Science	English	Average
0	61	24	81	55.333333
1	70	30	92	64.000000
2	96	84	84	88.000000
3	97	33	12	47.333333
4	31	62	11	34.666667
5	97	39	47	61.000000
6	11	73	69	51.000000
7	30	42	85	52.333333
8	67	31	98	65.333333
9	58	68	51	59.000000

DataFrame sorted by Average scores (descending order):

	Math	Science	English	Average
2	96	84	84	88.000000
8	67	31	98	65.333333
1	70	30	92	64.000000
5	97	39	47	61.000000
9	58	68	51	59.000000
0	61	24	81	55.333333
7	30	42	85	52.333333
6	11	73	69	51.000000
3	97	33	12	47.333333
4	31	62	11	34.666667

Filtered DataFrame with Average score > 50:

	Math	Science	English	Average
0	61	24	81	55.333333
1	70	30	92	64.000000
2	96	84	84	88.000000
5	97	39	47	61.000000
6	11	73	69	51.000000
7	30	42	85	52.333333
8	67	31	98	65.333333
9	58	68	51	59.000000

Statistics for Math scores:

Mean: 61.80, Standard Deviation: 30.36

Program 11. Write a Python program using NLTK to tokenize a sentence, filter out stopwords, and apply stemming?

12-Nov-2024

```
import nltk
from nltk.tokenize
import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
# Step 1: Download necessary NLTK data (run only once)
nltk.download('punkt')
nltk.download('stopwords')
# Step 2: Define a simple text
text = "Natural Language Processing with Python and NLTK is fun and interesting."
# Step 3: Tokenize the text into words
tokens = word_tokenize(text)
print("Tokenized words:", tokens)
# Step 4: Remove stopwords
stop_words = set(stopwords.words('english'))
filtered_tokens = [word for word in tokens if word.lower() not in stop_words]
print("\nFiltered words (stopwords removed):", filtered_tokens)
# Step 5: Perform stemming using PorterStemmer ps = PorterStemmer()
stemmed_tokens = [ps.stem(word) for word in filtered_tokens]
print("\nStemmed words:", stemmed_tokens)
```

```
Tokenized words: ['Natural', 'Language', 'Processing', 'with', 'Python', 'and', 'NLTK', 'is', 'fun', 'and',
'interesting', '.']
Filtered words (stopwords removed): ['Natural', 'Language', 'Processing', 'Python', 'NLTK', 'fun',
'interesting', '.']

Stemmed words: ['natur', 'languag', 'process', 'python', 'nltk', 'fun',
'interest', '.']
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
/root/nltk_data...[nltk_data] Package stopwords is
already up-to-date!
```

Program 12 . Implement a simple 3-layer neural network from scratch using Python and numpy. The network should have 2 input neurons, 2 hidden neurons, and 1 output neuron, using the sigmoid activation for both hidden and output layers. Train it on the XOR dataset using backpropagation for 100 epochs, printing the loss every 10 epochs. After training, print the predicted outputs for the XOR inputs. 15-Oct-2024

3- Layer Neural Network

Code 01

```
import numpy as np

# Sigmoid activation function
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

# Derivative of sigmoid for backpropagation
def sigmoid_derivative(x):
    return x * (1 - x)

# Neural Network class
class SimpleNN:

    def __init__(self, input_size, hidden_size, output_size):

        # Initialize weights randomly
        self.weights_input_hidden = np.random.rand(input_size, hidden_size)
        self.weights_hidden_output = np.random.rand(hidden_size, output_size)

        # Initialize biases
        self.bias_hidden = np.random.rand(hidden_size)
        self.bias_output = np.random.rand(output_size)

    # Forward pass
    def forward(self, inputs):
        # Input to hidden layer
        self.hidden_layer_activation = np.dot(inputs, self.weights_input_hidden) +
        self.bias_hidden
        self.hidden_layer_output = sigmoid(self.hidden_layer_activation)

        # Hidden to output layer
        self.output_layer_activation = np.dot(self.hidden_layer_output,
        self.weights_hidden_output) + self.bias_output
        self.output = sigmoid(self.output_layer_activation)
        return self.output

    # Backward pass (backpropagation)
    def backward(self, inputs, expected_output, learning_rate):
        # Calculate output error
        error_output = expected_output - self.output
        d_output = error_output * sigmoid_derivative(self.output)
        # Calculate hidden layer error
        error_hidden = d_output.dot(self.weights_hidden_output.T)
        d_hidden = error_hidden * sigmoid_derivative(self.hidden_layer_output)
        # Update weights and biases using gradient descent
        self.weights_hidden_output += self.hidden_layer_output.T.dot(d_output) * learning_rate
        self.bias_output += np.sum(d_output, axis=0) * learning_rate
        self.weights_input_hidden += inputs.T.dot(d_hidden) * learning_rate
        self.bias_hidden += np.sum(d_hidden, axis=0) * learning_rate

    # Train the model
    def train(self, inputs, expected_output, epochs, learning_rate):
        for epoch in range(epochs):
```

```

        # Forward pass
        self.forward(inputs)
        # Backward pass and weights update
        self.backward(inputs, expected_output, learning_rate) # Print loss every 100
        epochs
        if epoch % 100 == 0:
            loss = np.mean(np.square(expected_output - self.output))
            print(f"Epoch {epoch} - Loss: {loss}")

# XOR Dataset (example binary classification task)
inputs = np.array([[0, 0], [0, 1], [1, 0], [1, 1]]) # XOR inputs
expected_output = np.array([[0], [1], [1], [0]]) # XOR outputs

# Hyperparameters
input_size = 2 # Number of features in the input layer
hidden_size = 2 # Number of neurons in the hidden layer
output_size = 1 # Binary classification (1 neuron in the output layer)
epochs = 10000 # Number of training iterations
learning_rate = 0.1 # Learning rate for gradient descent

# Initialize the neural network

nn = SimpleNN(input_size, hidden_size, output_size)

# Train the neural network
nn.train(inputs, expected_output, epochs, learning_rate)

# Test the neural network
print("\nFinal Output after training:")
for i in range(len(inputs)):
    print(f"Input: {inputs[i]} => Predicted Output: {nn.forward(inputs[i])}")

Epoch 4800 - Loss: 0.2389717293167135
Epoch 4900 - Loss: 0.23650996180684464
Epoch 5000 - Loss: 0.23356503744226392
Epoch 5100 - Loss: 0.23009249514957203
Epoch 5200 - Loss: 0.2260781973462792
Epoch 5300 - Loss: 0.22155499126019743
Epoch 5400 - Loss: 0.21661067519423138
Epoch 5500 - Loss: 0.2113785955489492
Epoch 5600 - Loss: 0.2060102728930146
Epoch 5700 - Loss: 0.20064011835383688
Epoch 5800 - Loss: 0.19535597447337405
Epoch 5900 - Loss: 0.19018331330130778
Epoch 6000 - Loss: 0.1850820092745647
Epoch 6100 - Loss: 0.17994962284243815
Epoch 6200 - Loss: 0.17462560905954277
Epoch 6300 - Loss: 0.16889477834812883
Epoch 6400 - Loss: 0.16249369706971994
Epoch 6500 - Loss: 0.15512935714774995
Epoch 6600 - Loss: 0.14652318785381507
Epoch 6700 - Loss: 0.13649014855242383
Epoch 6800 - Loss: 0.12504485230934304
Epoch 6900 - Loss: 0.11249186228359834
Epoch 7000 - Loss: 0.09942605154513279
Epoch 7100 - Loss: 0.08659494845485854
Epoch 7200 - Loss: 0.07467784538617003
Epoch 7300 - Loss: 0.06411466613377548
Epoch 7400 - Loss: 0.055067453776417585
Epoch 7500 - Loss: 0.047486742293285034
Epoch 7600 - Loss: 0.04120794900431174
Epoch 7700 - Loss: 0.03602778936440759
Epoch 7800 - Loss: 0.031748569867997664
Epoch 7900 - Loss: 0.028197555053161016
Epoch 8000 - Loss: 0.025231903609445383
Epoch 8100 - Loss: 0.022736850694123265
Epoch 8200 - Loss: 0.02062147710636897
Epoch 8300 - Loss: 0.01881413809769125
Epoch 8400 - Loss: 0.01725838076982395
Epoch 8500 - Loss: 0.015909583012417905
Epoch 8600 - Loss: 0.014732295735133884
Epoch 8700 - Loss: 0.013698183717287711
Epoch 8800 - Loss: 0.012784445427901359
Epoch 8900 - Loss: 0.011972604280204108
Epoch 9000 - Loss: 0.011247583392884699

```

Epoch 9100 - Loss: 0.010596995188292924
Epoch 9200 - Loss: 0.01001059345668271
Epoch 9300 - Loss: 0.009479848422769856
Epoch 9400 - Loss: 0.008997615231420075
Epoch 9500 - Loss: 0.008557873699201503
Epoch 9600 - Loss: 0.008155522716156187
Epoch 9700 - Loss: 0.0077862167947832335
Epoch 9800 - Loss: 0.007446235316958131
Epoch 9900 - Loss: 0.007132377301425352

Final Output after training:

Input: [0 0] => Predicted Output: [0.08340633]
Input: [0 1] => Predicted Output: [0.92038212]
Input: [1 0] => Predicted Output: [0.92036973]
Input: [1 1] => Predicted Output: [0.08792555]

Program 13 . Write a Python program using Keras to build a deep learning model with 4 layers for binary classification. The model should have 2 hidden layers with 32 and 16 neurons respectively, each using the ReLU activation function. The output layer should use a sigmoid activation function. Use the XOR dataset as input, and train the model for 50 epochs. After training, print the final accuracy and the predicted output for each input. (1 point)

15-Oct-2024

4- Layer Deep Learning

Code 02

```
import numpy as np
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam

# Create a simple dataset (e.g., binary classification)
# Example: XOR dataset
X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]]) # Inputs
y = np.array([[0], [1], [1], [0]]) # Outputs

# Define the model
model = Sequential()

# Input layer (2 input neurons, corresponding to 2 features in X)
# Hidden layer 1: 64 neurons, activation: ReLU
model.add(Dense(64, input_dim=2, activation='relu'))

# Hidden layer 2: 32 neurons, activation: ReLU
model.add(Dense(32, activation='relu'))

# Hidden layer 3: 16 neurons, activation: ReLU
model.add(Dense(16, activation='relu'))

# Hidden layer 4: 8 neurons, activation: ReLU
model.add(Dense(8, activation='relu'))



# Output layer: 1 neuron (binary classification)
model.add(Dense(1, activation='sigmoid'))

# Compile the model
model.compile(optimizer=Adam(learning_rate=0.01), loss='binary_crossentropy', metrics=['accuracy'])

# Train the model
model.fit(X, y, epochs=1000, verbose=0)

# Evaluate the model performance
loss, accuracy = model.evaluate(X, y, verbose=0)
print(f'Final Accuracy: {accuracy * 100:.2f}%')

# Make predictions
predictions = model.predict(X)
print("\nPredictions:")
for i in range(len(X)):
    print(f"Input: {X[i]}, Predicted Output: {predictions[i][0]:.4f}")
```

 Final Accuracy: 100.00%
1/1  0s 65ms/step

Predictions:
Input: [0 0], Predicted Output: 0.0000
Input: [0 1], Predicted Output: 1.0000
Input: [1 0], Predicted Output: 1.0000
Input: [1 1], Predicted Output: 0.0000

Program 14. Setup single node Hadoop cluster and Apply HDFS Commands on Single Node Hadoop Cluster

22-Oct-2024

Setup single node Hadoop cluster and Apply HDFS Commands on Single Node Hadoop Cluster.

➤ Prerequisites

1. Java 8 runtime environment (JRE)
2. Apache Hadoop 3.3.6

Download Hadoop binary package

The first step is to download Hadoop binaries from the official website(<https://hadoop.apache.org/releases.html>). The binary package size is about 696 MB.

Unpack the package

After finishing the file download, we should unpack the package using 7zip or WinRAR.

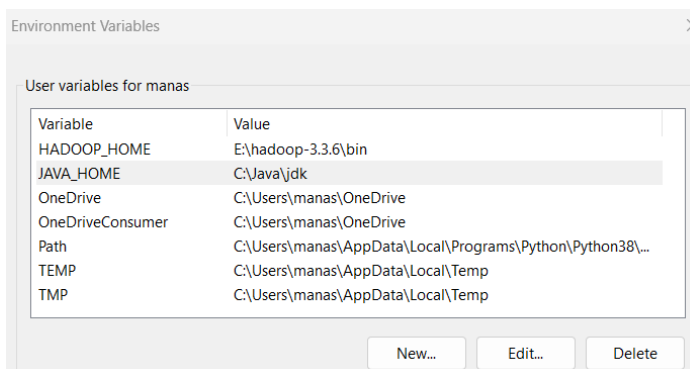
Java installation

Java is required to run Hadoop. If java is not installed, we have to install it. After finishing the file download we open a new command prompt, we should unpack the package.

Configure environment variables

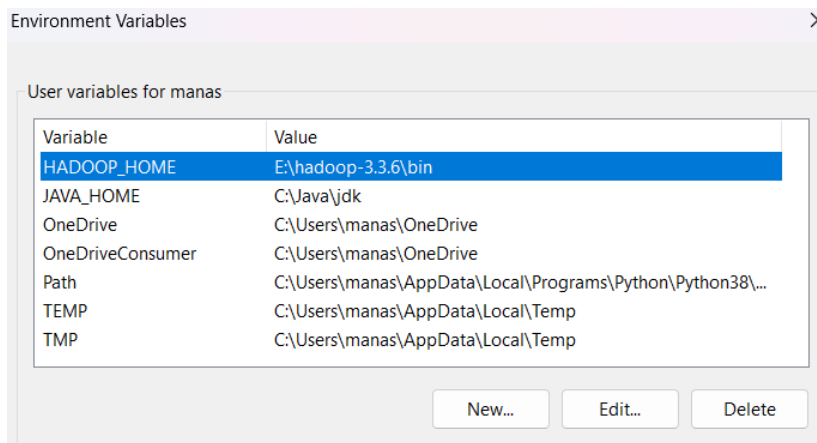
Now we've downloaded and unpacked all the files we need to configure two important environment variables.

We configure **JAVA_HOME** environment variable by adding new environment variable. Variable name : JAVA_HOME Variable value: Java path



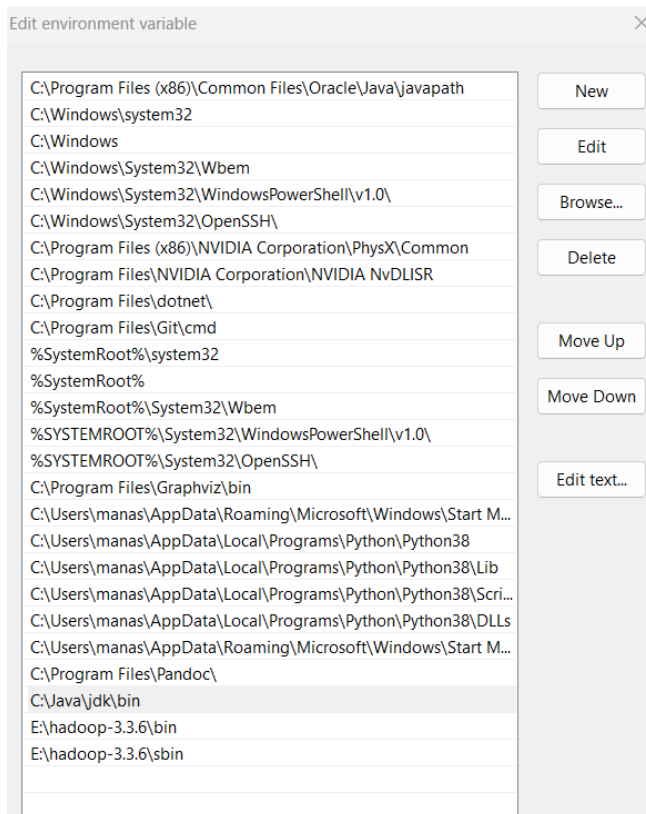
The same with **HADOOP_HOME** environment variable-

Variable name : HADOOP_HOME Variable value: C:\Hadoop\hadoop-3.3.6



- **Configure PATH environment variable**

Once we finish setting up the above two environment variables, we need to add the bin folders to the PATH environment variable.



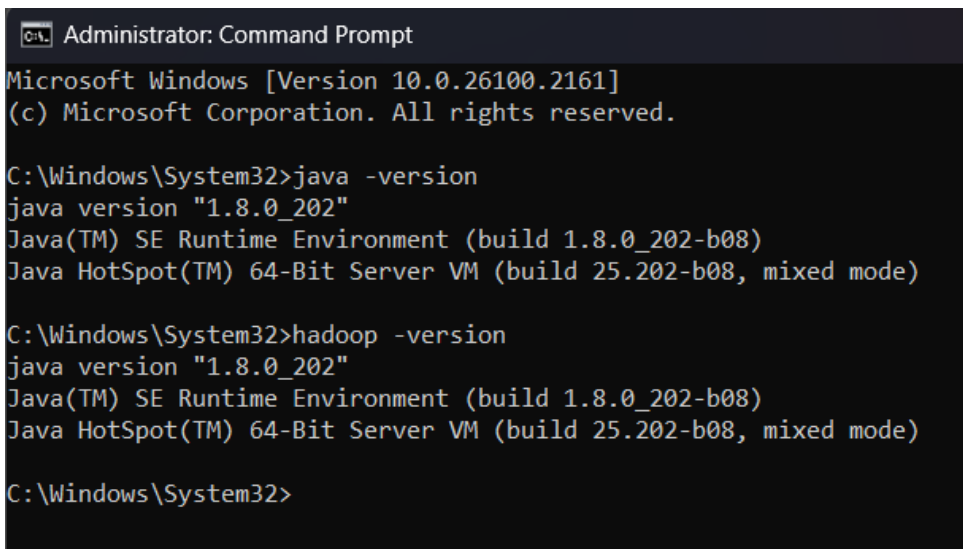
Verification of Installation

Once We complete the installation, Close terminal window and open a new one and run the following command to verify:

```
java -version
```

We can also be able to run the following command:

```
hadoop -version
```



```
Administrator: Command Prompt
Microsoft Windows [Version 10.0.26100.2161]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\System32>java -version
java version "1.8.0_202"
Java(TM) SE Runtime Environment (build 1.8.0_202-b08)
Java HotSpot(TM) 64-Bit Server VM (build 25.202-b08, mixed mode)

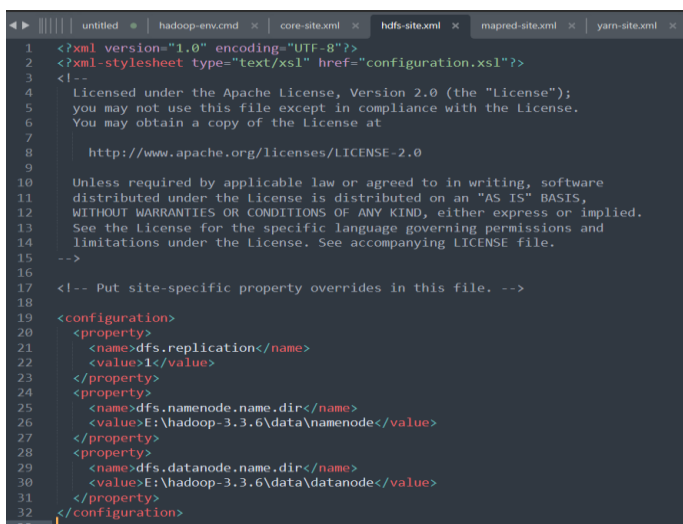
C:\Windows\System32>hadoop -version
java version "1.8.0_202"
Java(TM) SE Runtime Environment (build 1.8.0_202-b08)
Java HotSpot(TM) 64-Bit Server VM (build 25.202-b08, mixed mode)

C:\Windows\System32>
```

Configure Hadoop

Now we are ready to configure the most important part - Hadoop configurations which involves Core, YARN, MapReduce, HDFS configuration

Initialize HDFS



```
untitled | hadoop-env.cmd | core-site.xml | hdfs-site.xml | mapred-site.xml | yarn-site.xml
1  <?xml version="1.0" encoding="UTF-8"?>
2  <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
3  <!--
4  Licensed under the Apache License, Version 2.0 (the "License");
5  you may not use this file except in compliance with the License.
6  You may obtain a copy of the License at
7
8     http://www.apache.org/licenses/LICENSE-2.0
9
10 Unless required by applicable law or agreed to in writing, software
11 distributed under the License is distributed on an "AS IS" BASIS,
12 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
13 See the license for the specific language governing permissions and
14 limitations under the License. See accompanying LICENSE file.
15 -->
16
17 <!-- Put site-specific property overrides in this file. -->
18
19 <configuration>
20   <property>
21     <name>dfs.replication</name>
22     <value>1</value>
23   </property>
24   <property>
25     <name>dfs.namenode.name.dir</name>
26     <value>E:\hadoop-3.3.6\data\namenode</value>
27   </property>
28   <property>
29     <name>dfs.datanode.name.dir</name>
30     <value>E:\hadoop-3.3.6\data\datanode</value>
31   </property>
32 </configuration>
33
```

Run the following command in Command Prompt

hdfs namenode -format


```
C:\Windows\System32>hdfs namenode -format
2024-10-29 10:14:15,312 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG: host = Manas/192.168.1.196
STARTUP_MSG: args = [-format]
STARTUP_MSG: version = 3.3.6
STARTUP_MSG: classpath = E:\hadoop-3.3.6\etc\hadoop;E:\hadoop-3.3.6\share\hadoop\common;E:\hadoop-3.3.6\share\hadoop\com
mon\lib\animal-sniffer-annotations-1.17.jar;E:\hadoop-3.3.6\share\hadoop\common\lib\audience-annotations-0.5.0.jar;E:\h
adoop-3.3.6\share\hadoop\common\lib\avro-1.7.7.jar;E:\hadoop-3.3.6\share\hadoop\common\lib\checker-qual-2.5.2.jar;E:\ha
adoop-3.3.6\share\hadoop\common\lib\commons-beanutils-1.9.4.jar;E:\hadoop-3.3.6\share\hadoop\common\lib\commons-cli-1.2.j
ar;E:\hadoop-3.3.6\share\hadoop\common\lib\commons-codec-1.15.jar;E:\hadoop-3.3.6\share\hadoop\common\lib\commons-collec
tions-3.2.2.jar;E:\hadoop-3.3.6\share\hadoop\common\lib\commons-compress-1.21.jar;E:\hadoop-3.3.6\share\hadoop\common\li
b\commons-configuration2-2.8.0.jar;E:\hadoop-3.3.6\share\hadoop\common\lib\commons-daemon-1.0.13.jar;E:\hadoop-3.3.6\sha
re\hadoop\common\lib\commons-io-2.8.0.jar;E:\hadoop-3.3.6\share\hadoop\common\lib\commons-lang3-3.12.0.jar;E:\hadoop-3.3
.6\share\hadoop\common\lib\commons-logging-1.1.3.jar;E:\hadoop-3.3.6\share\hadoop\common\lib\commons-math3-3.1.1.jar;E:\
hadoop-3.3.6\share\hadoop\common\lib\commons-net-3.9.0.jar;E:\hadoop-3.3.6\share\hadoop\common\lib\commons-text-1.10.0.j
```

Start HDFS daemons

Run the following command to start HDFS daemons in Command Prompt:

start-all.cmd

Verify HDFS web portal UI through this link <http://localhost:9870/dfshealth.html#tab-overview>.



Overview 'localhost:9000' (active)


Started:	Tue Oct 29 10:27:26 +0530 2024
Version:	3.3.6, r1be78238728da9266a4f88195058f08fd012bf9c
Compiled:	Sun Jun 18 13:52:00 +0530 2023 by ubuntu from (HEAD detached at release-3.3.6-RC1)
Cluster ID:	CID-891bd78f-c90c-4979-b240-f46069696303
Block Pool ID:	BP-932695391-192.168.1.196-1730177794675



Datanode Information

✓ In service ⬇ Down ⚙ Decommissioning ⚠ Decommissioned ⛔ Decommissioned & dead
🟢 Entering Maintenance 🟡 In Maintenance 🔴 In Maintenance & dead

Datanode usage histogram



In operation

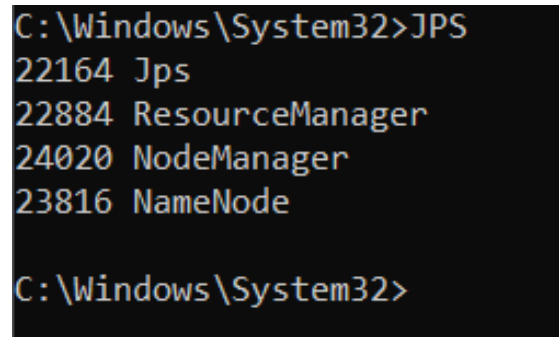
Start YARN daemons

Run the following command in an elevated Command Prompt window (Run as administrator) to start YARN daemons:

```
%HADOOP_HOME%\sbin\start-yarn.cmd
```

You can verify YARN resource manager UI when all services are started successfully.

<http://localhost:8088>




```
C:\Windows\System32>jps
22164 Jps
22884 ResourceManager
24020 NodeManager
23816 NameNode

C:\Windows\System32>
```


Program 15. Write a Python program to implement a Naive Bayes classifier for text classification using a dataset with text and label columns. Use CountVectorizer to preprocess the text data, train the model on a training set, and evaluate it on a test set with accuracy and F1-score metrics.

12-Nov-2024

```
from google.colab import drive
drive.mount('/content/drive')
```

 Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, f1_score
# Load the dataset (ensure it has 'text' and 'label' columns)
# You can replace this with the path to your dataset file
df = pd.read_csv('/content/drive/MyDrive/df_file.csv') # Ensure 'text' and 'label' columns
# Split the dataset into features (X) and labels (y)
X = df['text'] # Assuming the text data is in the 'text' column
y = df['label']
# Assuming the labels are in the 'label' column
# Split the data into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Initialize CountVectorizer for text preprocessing
vectorizer = CountVectorizer(stop_words='english')
# Fit the vectorizer on the training data and transform both training and testing data
X_train_vec = vectorizer.fit_transform(X_train)
X_test_vec = vectorizer.transform(X_test)
# Initialize and train the Naive Bayes classifier
nb_classifier = MultinomialNB()
nb_classifier.fit(X_train_vec, y_train)
# Predict on the test set
y_pred = nb_classifier.predict(X_test_vec)
# Evaluate the model using accuracy and F1-score
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
# Use 'weighted' for multi-class classification
# Print the results
print(f"Accuracy: {accuracy * 100:.2f}%")
print(f"F1-Score (Weighted): {f1:.2f}")
```

 Accuracy: 97.30%
F1-Score (Weighted): 0.97

Program 16. Write a Python program to calculate the Information Gain (IG) for a dataset and determine the root node for building a decision tree. The dataset contains attributes and a target column with binary values (e.g., "yes" or "no"). SEE ATTACHMENT to this Assignment.

19-Nov-2024

Outlook	Temperature	Humidity	Windy	Play Golf
Rainy	Hot	High	FALSE	No
Rainy	Hot	High	TRUE	No
Overcast	Hot	High	FALSE	Yes
Sunny	Mild	High	FALSE	Yes
Sunny	Cool	Normal	FALSE	Yes
Sunny	Cool	Normal	TRUE	No
Overcast	Cool	Normal	TRUE	Yes
Rainy	Mild	High	FALSE	No
Rainy	Cool	Normal	FALSE	Yes
Sunny	Mild	Normal	FALSE	Yes
Rainy	Mild	Normal	TRUE	Yes
Overcast	Mild	High	TRUE	Yes
Overcast	Hot	Normal	FALSE	Yes
Sunny	Mild	High	TRUE	No

```
import numpy as np
import pandas as pd
ANAND = pd.read_csv('/content/drive/MyDrive/golf-dataset.csv')

ANAND.columns

➡ Index(['Outlook', 'Temp', 'Humidity', 'Windy', 'Play Golf'], dtype='object')

def entropy(target_col):
    elements, counts = np.unique(target_col, return_counts=True)
    entropy = np.sum([(-counts[i]/np.sum(counts)) * np.log2(counts[i]/np.sum(counts)) for i
    in range(len(elements))])return entropy

def information_gain(data, split_attribute_name,
    target_name="class"):# Calculate the entropy of the
    total dataset
    total_entropy = entropy(data[target_name])

    # Calculate the values and the corresponding counts for
    the split attributevals, counts =
    np.unique(data[split_attribute_name], return_counts=True)

    # Calculate the weighted entropy
    weighted_entropy = np.sum([(counts[i]/np.sum(counts)) * entropy(data.where(data[split_attribute_name] ==
    vals[i]).dropna()[target_name])

    # Calculate the information gain
    information_gain = total_entropy -
    weighted_entropyreturn information_gain
```

```

print("Information Gain for 'Outlook':", information_gain(ANAND, 'Outlook', 'Play Golf'))
↔ Information Gain for 'Outlook': 0.24674981977443933

print("Information Gain for 'Temp':", information_gain(ANAND, 'Temp', 'Play Golf'))
↔ Information Gain for 'Temp': 0.02922256565895487

print("Information Gain for 'Humidity':", information_gain(ANAND, 'Humidity', 'Play Golf'))
↔ Information Gain for 'Humidity': 0.15183550136234159

print("Information Gain for 'Humidity':", information_gain(ANAND,
↔ 'Windy', 'Play Golf')) Information Gain for 'Humidity':
0.04812703040826949

```